

Web Recommender Systems 2025: Project

Davide Marchi
lnx547@alumni.ku.dk

Abstract

This report presents the development of the project for the Web Recommender System course held at the University of Copenhagen.

1 Week 6 [Feb 3 - Feb 9]: Familiarize Yourself with the Datasets

This week focused on understanding the dataset and preparing it for recommendation tasks. The goal was to clean the data, explore its properties, and extract key statistics to gain insights into user-item interactions. This preprocessing step was essential to ensure data quality for later modeling.

1.1 Dataset Description and Import

The dataset used in this project is the “Musical Instruments” dataset from the Amazon Review data 2023 (5-core subset). The dataset consists of 14,000 item ratings from nearly 900 users on approximately 500 items. The data is split into training and test sets, with 80% of the data allocated for training (“train.tsv”) and 20% for testing (“test.tsv”).

The dataset was imported using the Python library `pandas` ([pandas development team, 2020](#)) and its `read_csv` function.

1.2 Data Cleaning and Preprocessing

To ensure the quality of the dataset, I performed the following preprocessing steps:

- Removal of missing values (entries with missing user ID, item ID, or rating).
- Deduplication by keeping the most recent rating when a user rated the same item multiple times.
- Ensuring that all users in the test set are also present in the training set.

In terms of dimensions the effect of the cleaning can be summarized in the [Table 1](#).

Dataset	Before Cleaning	After Cleaning
Train Set	11000	9913
Test Set	3000	1822

Table 1: Effect of the Data Cleaning Process

1.3 Exploratory Data Analysis (EDA) and Statistics

To gain insights into the dataset, I computed the following statistics:

- Distribution of ratings per user in [Figure 1](#)
- Distribution of ratings per item in [Figure 2](#)
- Overall rating value distribution in [Figure 3](#)




Figure 1: Distribution of Ratings per User




Figure 2: Distribution of Ratings per Item




Figure 3: Distribution of Ratings per Value

Also, to show the long-tail distribution of ratings, I plotted the number of ratings per item in Figure 4.




Figure 4: Long tail distribution of rating over all the items

1.4 Highly Rated Items

To identify popular items, I computed the number of ratings ≥ 3 for each item and reported the top 5 most highly rated items in Table 2.

Item ID	High Ratings Count	Average High Rating
B0BSGM6CQ9	153	4.784
B0BPJ4Q6FJ	153	4.804
B09857JRP2	141	4.809
B0BCK6L7S5	120	4.583
B0BTC9YJ2W	119	4.815

Table 2: Top 5 Most Highly Rated Items

1.5 Discussion

The dataset exhibits common characteristics of user-item interactions:

- There is a long-tail distribution where a small number of items receive most of the ratings, while many items have very few ratings.
- The majority of users have rated only a few items, which can lead to sparsity issues in collaborative filtering approaches.
- Certain items dominate the high-rated category, which may bias recommendation systems.

2 Week 7 [Feb 10 - Feb 16]: Collaborative Filtering Recommender System

This week focused on implementing collaborative filtering recommender systems. The goal was to build models that predict user preferences based on past ratings, compare different approaches, and optimize their performance through hyperparameter tuning. These models serve as the foundation for generating personalized recommendations.

2.1 Top Popular (TopPop) Recommender System

As a baseline, I first implemented the **Top Popular (TopPop) recommender system**. This model recommends the most popular items based on the number of “high” ratings (ratings ≥ 3) in the training set. I computed the frequency of high ratings per item and selected the top-k items as recommendations for all users.

While computationally inexpensive, this method lacks personalization since it does not consider individual user preferences.

2.2 Collaborative Filtering Models

I then used the Scikit-Surprise library (Hug, 2020) to be able to compare it with one neighborhood-based model and one latent factor model. The ones I chose are:

- **User-based k-Nearest Neighbors with Means (KNNWithMeans):** A neighborhood-based approach that identifies similar users based on their rating history.
- **Singular Value Decomposition (SVD):** A latent factor model that decomposes the user-item interaction matrix into lower-dimensional representations.

2.3 Hyperparameter Tuning and Evaluation

To optimize both models and find the best hyperparameters, I performed a Grid Search with **5-fold cross-validation** on the training set. The parameters I tuned included:

- **KNNWithMeans**
 - **k:** [5, 10, 20, 30, 40, 50, 70, 100]
 - **sim_options:**
 - * **name:** ['cosine', 'msd']
 - * **user_based:** [True]
- **SVD**
 - **n_factors:** [3, 5, 10, 20, 50, 100]
 - **n_epochs:** [5, 10, 20, 50, 100]

The evaluation metric chosen was **Root Mean Square Error (RMSE)**. RMSE penalizes larger errors more than smaller ones, making it a good measure to capture the overall predictive accuracy of our recommendation models. Lower RMSE indicates that the model predictions are closer to the actual ratings.

2.4 Results and Discussion

The best hyperparameters for each model can be found in [Table 3](#).

Model	Best Hyperparameters	Best RMSE
KNNWithMeans	'k': 70, 'sim_options': {'name': 'cosine', 'user_based': True}	0.9113
SVD	'n_factors': 10, 'n_epochs': 20	0.8445

Table 3: Best Hyperparameters and RMSE for Each Model

The SVD model, despite its low dimensional latent space, outperformed the KNNWithMeans model with a lower RMSE of 0.8445 compared to 0.9113. This indicates that the latent factor model is better at capturing the underlying patterns in the user-item interaction matrix.

After that both models have been trained again on the whole training set, to make sure that no data available for training is left behind before the evaluation on the test set.

3 Week 8 [Feb 17 - Feb 23]: Evaluation of Recommender Systems

This week focused on evaluating the performance of the collaborative filtering models implemented in Week 7. The evaluation aimed to measure both the accuracy of rating predictions and the effectiveness of recommendations.

3.1 RMSE Evaluation

The first step in evaluating the models was measuring the **Root Mean Square Error (RMSE)** on the test set. RMSE quantifies how close the predicted ratings are to the actual ratings, with lower values indicating better accuracy. The trained models were tested on the preprocessed test data, and their RMSE values were computed and are shown in [Table 4](#).

Model	Validation RMSE	Test RMSE
KNNWithMeans	0.9113	1.0619
SVD	0.8445	0.9796

Table 4: RMSE Comparison on Test Set

While RMSE is useful for evaluating rating prediction accuracy, it has some limitations. It does not consider ranking quality and does not directly reflect recommendation usefulness. For this reason, additional ranking-based metrics were computed.

3.2 Ranking-based Evaluation Metrics

To assess the quality of the generated recommendations, the models were evaluated using the following metrics:

- **Hit Rate:** Measures how often at least one relevant item appears in the top-k recommendations.
 - **Advantages:** Simple, easy to interpret.
 - **Disadvantages:** Doesn't consider ranking or multiple relevant items.
- **Precision@k:** Proportion of recommended items in the top-k list that are relevant.
 - **Advantages:** Direct measure of recommendation quality.
 - **Disadvantages:** Ignores total relevant items available.
- **MAP@k:** Averages precision at positions of relevant items, rewarding good ranking.
 - **Advantages:** Considers ranking order.

- **Disadvantages:** Sensitive to users with few relevant items.
- **MRR@k:** Focuses on the rank of the first relevant item. Computed as one divided by its rank.
 - **Advantages:** Rewards early relevant recommendations.
 - **Disadvantages:** Ignores additional relevant items.
- **Coverage:** Proportion of unique items recommended at least once.
 - **Advantages:** Indicates diversity in recommendations.
 - **Disadvantages:** High coverage doesn't mean better relevance.

The results of the computation of these metrics are shown in [Table 5](#).

Model	Hit Rate	Precision@10	MAP@10	MRR@10	Coverage
KNNWithMeans	0.1092	0.0117	0.0084	0.0314	0.5874
SVD	0.0850	0.0095	0.0066	0.0281	0.1552
TopPop	0.2573	0.0325	0.0339	0.1187	0.0196

Table 5: Evaluation Metrics for Top-10 Recommendations

3.3 Results and Discussion

- **Hit Rate:** Surprisingly, the **TopPop model** achieved the highest hit rate, likely because popular items frequently appear in the test set.
- **Precision@10:** The TopPop model outperformed the collaborative filtering methods in precision as well, suggesting that highly rated items tend to be good general recommendations.
- **MAP@10 and MRR@10:** TopPop also performed best in these ranking-based metrics, showing that the most popular items are often relevant.
- **Coverage:** The **kNN model** had the highest coverage (58.7%), indicating that it provided more diverse recommendations, while SVD had lower coverage. The TopPop model had the lowest coverage, meaning it recommended only a small subset of items.

The evaluation results show an interesting (but expected) trade-off: The **TopPop model** is simple but performs surprisingly well, especially in terms

of hit rate and ranking metrics. However, it lacks personalization. The **kNN model** achieves the highest coverage, meaning it suggests a wider variety of items, which could be useful for expanding user interest. The **SVD model** despite achieving the lowest RMSE, performed worse than expected in ranking-based evaluation, possibly due to the sparsity of the dataset.

These results suggest that, while collaborative filtering methods are typically expected to perform well, the characteristics of the dataset (such as sparsity and rating distribution) play a significant role in determining effectiveness.

In the end we do not have a clear winner, as the best model depends on the specific use case and the trade-offs between accuracy, diversity, and coverage.

References

- Nicolas Hug. 2020. [Surprise: A python library for recommender systems](#). *Journal of Open Source Software*, 5(52):2174.
- The pandas development team. 2020. [pandas-dev/pandas: Pandas](#).