

Web Recommender Systems 2025: Project

Davide Marchi
lnx547@alumni.ku.dk

1 Week 6 [Feb 3 - Feb 9]: Familiarize Yourself with the Datasets

This week focused on understanding the dataset and preparing it for recommendation tasks. The goal was to clean the data, explore its properties, and extract key statistics to gain insights into user-item interactions. This preprocessing step was essential to ensure data quality for later modeling.

1.1 Dataset Description

The dataset used in this project is the “Musical Instruments” dataset from the Amazon Review data 2023 (5-core subset). The dataset consists of 14,000 item ratings from nearly 900 users on approximately 500 items. The data is split into training and test sets, with 80% of the data allocated for training (“train.tsv”) and 20% for testing (“test.tsv”).

1.2 Data Cleaning and Preprocessing

To ensure the quality of the dataset, I performed the following preprocessing steps:

1. Removal of missing values (entries with missing user ID, item ID, or rating).
Outcome: 0 ratings removed.
2. Deduplication by keeping the most recent rating when a user rated the same item multiple times.
Outcome: 1087 ratings removed from the training set and 1178 from test set.
3. Ensuring that all users in the test set are also present in the training set.
Outcome: 0 users removed.

In terms of dimensions the effect of the cleaning can be summarized in [Table 1](#) (for the ratings) and [Table 2](#) (for the users).

Dataset	Before Cleaning	After Cleaning
Train Set	11000	9913
Test Set	3000	1822

Table 1: Effect of the Data Cleaning Process on the Number of Ratings

Dataset	Before Cleaning	After Cleaning
Train Set	800	800
Test Set	437	437

Table 2: Effect of the Data Cleaning Process on the Number of Users

1.3 Exploratory Data Analysis and Statistics

To gain insights into the dataset, I computed the following statistics **on the training set**:

- Distribution of ratings per user in [Figure 1](#)
- Distribution of ratings per item in [Figure 2](#)
- Overall rating value distribution in [Figure 3](#)




Figure 1: Distribution of Ratings per User




Figure 2: Distribution of Ratings per Item




Figure 3: Distribution of Ratings per Value

Also, to show the long-tail distribution of ratings, I plotted the number of ratings per item in Figure 4, computed on the training set.




Figure 4: Long tail distribution of rating over all the items

1.4 Highly Rated Items

To identify popular items, I computed the number of ratings ≥ 3 for each item and reported the top 5 most highly rated items in Table 3.

Item ID	High Ratings	Average Rating
B0BSGM6CQ9	153	4.6899
B0BPJ4Q6FJ	153	4.7372
B09857JRP2	141	4.7692
B0BCK6L7S5	120	4.3030
B0BTC9YJ2W	119	4.7377

Table 3: Top 5 Most Highly Rated Items

1.5 Discussion

The dataset exhibits common characteristics of user-item interactions:

- There is a long-tail distribution where a small number of items receive most of the ratings, while many items have very few ratings.
- The majority of users have rated only a very small percentage of the available items, which can lead to sparsity issues in collaborative filtering approaches. Indeed, the user-item matrix will be mostly empty, so there are often very few overlapping ratings between users. This makes it harder for algorithms like kNN to find similar users, and for models like SVD to learn meaningful patterns, since these only emerge clearly when there's enough data to separate them from noise.
- Few items dominate the high-rated category, which may bias recommendation systems and favor a TopPop approach.
- Most ratings are 5/5, likely because users tend to rate items they like more than those they dislike. This results in a highly imbalanced dataset for the task of predicting ratings.

2 Week 7 [Feb 10 - Feb 16]: Collaborative Filtering Recommender System

This week focused on implementing collaborative filtering recommender systems. The goal was to build models that predict user preferences based on past ratings, compare different approaches, and optimize their performance through hyperparameter tuning. These models serve as the foundation for generating personalized recommendations.

2.1 Top Popular (TopPop) Recommender System

As a baseline, I first implemented the **Top Popular (TopPop) recommender system**. This model recommends the most popular items based on the

number of “high” ratings (ratings ≥ 3) in the training set. I computed the frequency of high ratings per item and selected the top-k items as recommendations for all users.

While computationally inexpensive, this method lacks personalization since it does not consider individual user preferences.

2.2 Collaborative Filtering Models

I then used the Scikit-Surprise library(Hug, 2020) to be able to compare TopPop with one neighborhood-based model and one latent factor model. The ones I chose are:

- **User-based k-Nearest Neighbors with Means (KNNWithMeans):** A neighborhood-based approach that identifies similar users based on their rating history.
- **Singular Value Decomposition (SVD):** A latent factor model that decomposes the user-item interaction matrix into lower-dimensional representations.

2.3 Hyperparameter Tuning and Evaluation

To optimize both models and find the best hyperparameters, I performed a Grid Search with **5-fold cross-validation** on the training set. The parameters I tuned included:

- **User-based KNNWithMeans**
 - **k:** [5, 10, 20, 30, 40, 50, 70, 100]
 - **similarity_measure:** ['cosine', 'msd']
 - * **cosine:** Computes the cosine of the angle between two vectors in a multi-dimensional space, ignoring magnitude and focusing on the direction of ratings.
 - * **msd:** Computes the Mean Squared Difference between corresponding elements of two vectors, capturing the absolute differences in rating values.
- **SVD**
 - **n_factors:** [3, 5, 10, 20, 50, 100]
 - **n_epochs:** [5, 10, 20, 50, 100]

The evaluation metric chosen was **Root Mean Square Error (RMSE)**. RMSE penalizes larger errors more than smaller ones, making it a good measure to capture the overall predictive accuracy of our recommendation models. Lower RMSE indicates that the model predictions are closer to the actual ratings.

2.4 Results and Discussion

The best hyperparameters for each model can be found in Table 4.

Model	Best Hyperparameters	Avg RMSE
KNNWithMeans	'k': 70, 'similarity_measure': 'cosine'	0.9100
SVD	'n_factors': 5, 'n_epochs': 20	0.8440

Table 4: Best Hyperparameters and Average RMSE Across the 5 Folds for Each Model

The SVD model, despite its low dimensional latent space, outperformed the KNNWithMeans model with a lower RMSE of 0.8440 compared to 0.9100. This indicates that the latent factor model is better at capturing the underlying patterns in the user-item interaction matrix.

After that, I ran the models with the optimal hyperparameters on the full training set and used them to predict the missing ratings. These predictions will later be compared to the actual ratings in the test set to assess model performance.

3 Week 8 [Feb 17 - Feb 23]: Evaluation of Recommender Systems

This week focused on evaluating the performance of the collaborative filtering models implemented in Week 7. The evaluation aimed to measure both the accuracy of rating predictions and the effectiveness of recommendations.

3.1 RMSE Evaluation

The first step in evaluating the models was measuring the **Root Mean Square Error (RMSE)** on the test set. RMSE quantifies how close the predicted ratings are to the actual ratings.

The trained models were tested on the preprocessed test data, and their RMSE values were computed and are shown in Table 5.

Model	Validation RMSE	Test RMSE
KNNWithMeans	0.9100	1.0619
SVD	0.8440	0.9795

Table 5: RMSE Comparison on Test Set

While RMSE is useful for evaluating rating prediction accuracy, it has some limitations. It does not consider ranking quality and does not directly reflect recommendation usefulness. For this reason, additional ranking-based metrics were computed.

3.2 Ranking-based Evaluation Metrics

To assess the quality of the generated recommendations, the models were evaluated **on the test set** using the following metrics:

- **Hit Rate:** Measures how often at least one relevant item appears in the top-k recommendations.
 - **Advantages:** Simple, easy to interpret.
 - **Disadvantages:** Doesn't consider ranking or multiple relevant items.
- **Precision@k:** Proportion of recommended items in the top-k list that are relevant.
 - **Advantages:** Direct measure of recommendation quality.
 - **Disadvantages:** Ignores total relevant items available and ranking order.
- **MAP@k:** Averages precision at positions of relevant items, rewarding good ranking.
 - **Advantages:** Considers ranking order.
 - **Disadvantages:** Sensitive to users with few relevant items.
- **MRR@k:** Focuses on the rank of the first relevant item. Computed as one divided by its rank.
 - **Advantages:** Rewards early relevant recommendations.
 - **Disadvantages:** Ignores additional relevant items.
- **Coverage:** Proportion of unique items recommended at least once.
 - **Advantages:** Indicates diversity in recommendations.
 - **Disadvantages:** High coverage doesn't mean better relevance.

The results of the computation of these metrics are shown in [Table 6](#).

Model	Hit Rate	Precision@10	MAP@10	MRR@10	Coverage
KNNWithMeans	0.1092	0.0117	0.0084	0.0314	0.5874
SVD	0.0825	0.0083	0.0061	0.0222	0.1179
TopPop	0.2573	0.0325	0.0339	0.1187	0.0196

Table 6: Evaluation Metrics for Top-10 Recommendations **on the Test Set**

3.3 Results and Discussion

- **Hit Rate:** Unsurprisingly, the **TopPop model** achieved the highest hit rate, likely because popular items frequently appear in the test set.
- **Precision@10:** The TopPop model outperformed the collaborative filtering methods in precision as well, suggesting that highly rated items tend to be good general recommendations.
- **MAP@10 and MRR@10:** TopPop also performed best in these ranking-based metrics, showing that the most popular items are often relevant.
- **Coverage:** The **kNN model** had the highest coverage (58.7%), indicating that it provided more diverse recommendations, while SVD had lower coverage. The TopPop model had the lowest coverage, meaning it recommended only a small subset of items.

The evaluation results show an interesting (but expected) trade-off:

- The **TopPop model** is simple but performs surprisingly well, especially in terms of hit rate and ranking metrics. However, it completely lacks of any sort of personalization.
- The **kNN model** achieves the highest coverage, meaning it suggests a wider variety of items, which could be useful for expanding user interest.
- The **SVD model** despite achieving the lowest RMSE, performed worse than expected in ranking-based evaluation, possibly due to the sparsity of the dataset.

3.4 What can be concluded

These results suggest that, while collaborative filtering methods are typically expected to perform well, the characteristics of the dataset (such as sparsity and rating distribution) play a significant role in determining effectiveness.

At this point, there is no clear winner, as the best model depends on the specific use case and the **trade-offs** between accuracy and coverage.

In a case where the only important thing is precision, **TopPop** is undoubtedly the best choice here. But the real hard yet rewarding thing to obtain is a wide coverage capable of recommending (even

if with less confidence than TopPop) some niche items reviewed just by few users. In this regard, the higher coverage of **KNNWithMeans** proves to be very useful.

This approach also helps to spread the upcoming reviews, while **TopPop** would just end up making the popular items even more popular and the less common items even less common in proportion.

3.5 Analysis on the Effect of the Long-tail

To better understand how user and item activity impacts recommendation quality, I analyzed the performance of the models across different segments of the user and item distribution.

Users were sorted by the number of ratings in the training set and divided into top 20% and bottom 20% groups. Then, I computed on the test set the Hit Rate for each group using the three recommenders (TopPop, KNNWithMeans, SVD). The results are reported in Table 7.

Model	Top 20% Users	Bottom 20% Users
TopPop	0.164	0.371
KNNWithMeans	0.180	0.145
SVD	0.098	0.119

Table 7: Hit Rate for Top and Bottom 20% of Users

Interestingly, TopPop performs best for less active users. This aligns with its non-personalized nature: popular items are more likely to match the sparse preferences of cold-start users. On the other hand, KNN performs better for active users, benefiting from richer collaborative signals.

A similar analysis was performed for items, by computing on the test set the coverage for the top and bottom 20% most rated items of the train set. Table 8 shows the results.

Model	Top 20% Items	Bottom 20% Items
TopPop	0.099	0.000
KNNWithMeans	0.475	0.624
SVD	0.188	0.020

Table 8: Coverage for Top and Bottom 20% of Items

KNNWithMeans significantly outperforms other models in recommending less frequent items, highlighting its ability to diversify suggestions. In contrast, TopPop completely ignores the long tail.

3.6 Error Analysis for the Neighborhood-based CF

To better understand the behavior of KNNWithMeans, I performed an error analysis based on

Reciprocal Rank (RR). I selected one user with high RR (= 1.0) and one with low RR (= 0.0), and inspected their top-10 nearest neighbors.

I found that the average overlap in item history between the high-RR user and their neighbors was **1.30 items**, compared to only **1.00 item** for the low-RR user. This suggests that the model performs poorly when the user's rating history is dissimilar to their neighbors'.

Additionally, in cases with low RR, I noticed that many of the nearest neighbors had very different item preferences, which diluted the quality of the recommendations.

Conclusion: For KNN-based models to be effective, the training data must have sufficient overlap between users. Sparse data significantly impairs performance and increases the chance of irrelevant neighbors.

4 Week 9 [Feb 24 - Mar 2]: Text Representation

This week focused on Natural Language Processing (NLP) techniques to extract useful representations from the item descriptions. These textual features serve as the foundation for building content-based and hybrid recommendation models.

4.1 Preprocessing

To reduce noise and focus on semantic content, I cropped all item descriptions to the first 1000 characters, which covered more than 84% of items. Then, I applied the following preprocessing steps:

- Tokenization using NLTK
- Lowercasing
- Stopword removal (NLTK stopwords list)
- Stemming (Snowball Stemmer)

The vocabulary size was reduced from 51,296 to 27,243 tokens, significantly improving sparsity and reducing dimensionality.

4.2 Vector Representations

Two main representations were used:

- **TF-IDF:** Sparse vectors built from the preprocessed descriptions using Scikit-learn. Final shape: (matrix $27,243$).

- **Word2Vec**: Pretrained Google News embeddings were averaged over tokens to produce 300-dimensional dense vectors.

4.3 Cosine Similarity Analysis

I selected a subset of 10 items and compared their pairwise similarities using each method. Figure 5 shows the heatmaps.




Figure 5: Cosine similarity between item vectors (TF-IDF vs Word2Vec)

TF-IDF captures lexical overlap, while Word2Vec captures conceptual similarity, even when word choice varies. The two spaces complement each other well.

4.4 Optional: BERT Embeddings

To further explore semantic encoding, I used the [CLS] token representation from BERT. This produced 768-dimensional dense vectors. Similarity matrices from BERT embeddings showed better clustering for semantically close items, though the cost was significantly higher.

5 Week 10 [Mar 3 - Mar 9]: Content-Based Recommender System

Using the representations from Week 9, I implemented a recommender system that ranks items based on similarity to a user's content profile.

5.1 Item and User Representations

Each item was represented as a TF-IDF vector over both its title and description. The vectors were concatenated, resulting in a 41,550-dimensional feature space. Titles were preprocessed identically to descriptions.

Each user was represented as the mean vector of the items they rated in the training set.

5.2 Prediction and Evaluation

I used cosine similarity between user and item vectors to rank unobserved items. The predicted rating was a linear transformation of similarity to the [1, 5] scale.

Evaluation metrics, using ratings ≥ 4 in the test set, are presented in Table 9.

Model	Hit Rate	Precision@10	MAP@10	MRR@10	Coverage
Content-Based	0.0121	0.0012	0.0007	0.0022	0.0485

Table 9: Content-Based Recommender Evaluation Results

The model struggled to make strong predictions, likely due to the weak correlation between text similarity and rating preferences. However, it provided valuable diversity and cold-start handling capabilities.

6 Week 11 [Mar 10 - Mar 16]: Hybrid Recommender Systems

To combine the strengths of collaborative and content-based methods, I implemented three hybrid strategies.

6.1 Parallel Hybrid

Both models rank all items, and scores are aggregated by inverse rank sum. Results in Table 10.

6.2 Switching Hybrid

Users with ≥ 5 ratings use KNN, others use content-based.

6.3 Pipeline Hybrid

Content-based model generates candidates (top-25), then re-ranked using KNN.

Model	Hit Rate	Precision@10	MAP@10	MRR@10	Coverage
Parallel Hybrid	0.0583	0.0058	0.0048	0.0186	0.0944
Switching Hybrid	0.0825	0.0083	0.0077	0.0262	0.0526
Pipeline Hybrid	0.0194	0.0019	0.0031	0.0133	0.0486

Table 10: Evaluation Results of Hybrid Recommenders

6.4 Discussion

The switching hybrid performed best overall, combining personalization with fallback for cold-start users. The parallel hybrid showed the highest coverage improvement, proving beneficial for diversity. Pipeline hybrid had moderate impact, possibly limited by content-only candidate generation.

In conclusion, hybrid approaches allow for more robust and balanced recommendations by leveraging both collaborative signals and item content.

References

Nicolas Hug. 2020. [Surprise: A python library for recommender systems](#). *Journal of Open Source Software*, 5(52):2174.